

# Data Cleaning

Andrew Jaffe

July 11, 2017

# Data Cleaning

In general, data cleaning is a process of investigating your data for inaccuracies, or recoding it in a way that makes it more manageable.

**MOST IMPORTANT RULE - LOOK AT YOUR DATA!**

## Useful checking functions

- `is.na` - is TRUE if the data is FALSE otherwise
- `!` - negation (NOT)
  - if `is.na(x)` is TRUE, then `!is.na(x)` is FALSE
- `all` takes in a `logical` and will be TRUE if ALL are TRUE
  - `all(!is.na(x))` - are all values of x NOT NA
- `any` will be TRUE if ANY are true
  - `any(is.na(x))` - do we have any NA's in x?
- `complete.cases` - returns TRUE if EVERY value of a row is NOT NA
  - very stringent condition
  - FALSE missing one value (even if not important)

# Dealing with Missing Data

## Missing data types

One of the most important aspects of data cleaning is missing values.

Types of "missing" data:

- **NA** - general missing data
- **NaN** - stands for "**N**ot **a** **N**umber", happens when you do  $0/0$ .
- **Inf** and **-Inf** - Infinity, happens when you take a positive number (or negative number) by 0.

## Finding Missing data

Each missing data type has a function that returns **TRUE** if the data is missing:

- NA - `is.na`
- NaN - `is.nan`
- Inf and -Inf - `is.infinite`
- `is.finite` returns **FALSE** for all missing data and **TRUE** for non-missing

## Missing Data with Logicals

One important aspect (esp with subsetting) is that logical operations return **NA** for **NA** values. Think about it, the data could be  $> 2$  or not we don't know, so R says there is no **TRUE** or **FALSE**, so that is missing:

```
x = c(0, NA, 2, 3, 4)
x > 2
```

```
[1] FALSE    NA FALSE  TRUE  TRUE
```

## Missing Data with Logicals

What to do? What if we want if  $x > 2$  and  $x$  isn't NA?  
Don't do  $x \neq \text{NA}$ , do  $x > 2$  and  $x$  is NOT NA:

```
x != NA
```

```
[1] NA NA NA NA NA
```

```
x > 2 & !is.na(x)
```

```
[1] FALSE FALSE FALSE  TRUE  TRUE
```



## Missing Data with Logicals

What about seeing if a value is equal to multiple values? You can do `(x == 1 | x == 2) & !is.na(x)`, but that is not efficient.

```
(x == 0 | x == 2) # has NA
```

```
[1] TRUE NA TRUE FALSE FALSE
```

```
(x == 0 | x == 2) & !is.na(x) # No NA
```

```
[1] TRUE FALSE TRUE FALSE FALSE
```

what to do?

## Missing Data with Logicals: %in%

Introduce the %in% operator:

```
x %in% c(0, 2) # NEVER has NA and returns logical
```

```
[1] TRUE FALSE TRUE FALSE FALSE
```

reads "return TRUE if x is in 0 or 2". (Like `inlist` in Stata).

## Missing Data with Logicals: %in%

NEVER has NA, even if you put it there (BUT DON'T DO THIS):

```
x %in% c(0, 2, NA) # NEVER has NA and returns logical
```

```
[1] TRUE TRUE TRUE FALSE FALSE
```

```
x %in% c(0, 2) | is.na(x)
```

```
[1] TRUE TRUE TRUE FALSE FALSE
```

## Missing Data with Operations

Similarly with logicals, operations/arithmetic with **NA** will result in **NAs**:

```
x + 2
```

```
[1]  2 NA  4  5  6
```

```
x * 2
```

```
[1]  0 NA  4  6  8
```

# Tables and Tabulations

## Useful checking functions

- `unique` - gives you the unique values of a variable
- `table(x)` - will give a one-way table of `x`
  - `table(x, useNA = "ifany")` - will have row NA
- `table(x, y)` - will give a cross-tab of `x` and `y`

# Creating One-way Tables

Here we will use `table` to make tabulations of the data. Look at `?table` to see options for missing data.

```
unique(x)
```

```
[1] 0 NA 2 3 4
```

```
table(x)
```

```
x
0 2 3 4
1 1 1 1
```

```
table(x, useNA = "ifany") # will not
```

```
x
 0  2  3  4 <NA>
1  1  1  1  1  1
```

# Creating One-way Tables

`useNA = "ifany"` will not have NA in table heading if no NA:

```
table(c(0, 1, 2, 3, 2, 3, 3, 2,2, 3),  
      useNA = "ifany")
```

```
0 1 2 3  
1 1 4 4
```



# Creating One-way Tables

You can set `useNA = "always"` to have it always have a column for NA

```
table(c(0, 1, 2, 3, 2, 3, 3, 2,2, 3),  
      useNA = "always")
```

	0	1	2	3	<NA>
1	1	1	4	4	0

# Tables with Factors

If you use a **factor**, all levels will be given even if no exist! - (May be wanted or not):

```
fac = factor(c(0, 1, 2, 3, 2, 3, 3, 2,2, 3),  
             levels = 1:4)  
tab = table(fac)  
tab
```

```
fac  
1 2 3 4  
1 4 4 0
```

```
tab[ tab > 0 ]
```

```
fac  
1 2 3  
1 4 4
```

# Creating Two-way Tables

A two-way table. If you pass in 2 vectors, **table** creates a 2-dimensional table.

```
tab <- table(c(0, 1, 2, 3, 2, 3, 3, 2, 2, 3),  
             c(0, 1, 2, 3, 2, 3, 3, 4, 4, 3),  
             useNA = "always")
```

## Recoding to missing

Sometimes people code missing data in weird or inconsistent ways.

```
ages = c(23, 21, 44, 32, 57, 65, -999, 54)  
range(ages)
```

```
[1] -999    65
```

## Recoding to missing

How do we change the -999 to be treated as missing?

```
ages[ages == -999] = NA  
range(ages)
```

```
[1] NA NA
```

```
range(ages, na.rm=TRUE)
```

```
[1] 21 65
```

## Recoding from missing

What if you were the person that coded the -999

```
is.na(ages)
```

```
[1] FALSE FALSE FALSE FALSE FALSE FALSE TRUE FALSE
```

```
ages[is.na(ages)] = -999  
ages
```

```
[1] 23 21 44 32 57 65 -999 54
```

## Read in the UFO dataset

- Download data from [http://sisbid.github.io/Module1/data/ufo/ufo\\_data\\_complete.csv.gz](http://sisbid.github.io/Module1/data/ufo/ufo_data_complete.csv.gz)
- Extract the CSV from the zipped file
- Save it (or move it) to the same folder as your day1.R script

```
ufo = read_csv("../data/ufo/ufo_data_complete.csv")
```

Parsed with column specification:

```
cols(  
  datetime = col_character(),  
  city = col_character(),  
  state = col_character(),  
  country = col_character(),  
  shape = col_character(),  
  `duration (seconds)` = col_double(),  
  `duration (hours/min)` = col_character(),  
  comments = col_character(),  
  `date posted` = col_character(),  
  latitude = col_character(),  
  longitude = col_double()  
)
```

# Checking for logical conditions

- `any()` - checks if there are any TRUEs
- `all()` - checks if ALL are true

```
any(is.na(ufo$state)) # are there any NAs?
```

```
[1] TRUE
```

```
table(is.na(ufo$state)) # are there any NAs?
```

```
FALSE  TRUE  
81356  7519
```



# Recoding Variables

## Example of Recoding: base R

For example, let's say gender was coded as Male, M, m, Female, F, f. Using Excel to find all of these would be a matter of filtering and changing all by hand or using if statements.

In R, you can simply do something like:

```
data$gender[data$gender %in%  
  c("Male", "M", "m")] <- "Male"
```

## Example of Cleaning: more complicated

Sometimes though, it's not so simple. That's where functions that find patterns come in very useful.

```
table(gender)
```

```
gender
  F FeMAle FEMALE      Fm      M      Ma      mAle      Male      MaLe      MALE
75      82      74      89      89      79      87      89      88      95
Man  Woman
73      80
```

# String functions

# Useful String Functions

## Useful String functions

- `toupper()`, `tolower()` - uppercase or lowercase your data:
- `str_trim()` (in the `stringr` package) or `trimws` in base
  - will trim whitespace
- `nchar` - get the number of characters in a string
- `paste()` - paste strings together with a space
- `paste0` - paste strings together with no space as default

## Pasting strings with `paste` and `paste0`

Paste can be very useful for joining vectors together:

```
paste("Visit", 1:5, sep = "_")
```

```
[1] "Visit_1" "Visit_2" "Visit_3" "Visit_4" "Visit_5"
```

```
paste("Visit", 1:5, sep = "_", collapse = " ")
```

```
[1] "Visit_1 Visit_2 Visit_3 Visit_4 Visit_5"
```

```
paste("To", "is going be the ", "we go to the store!", sep = "day ")
```

```
[1] "Today is going be the day we go to the store!"
```

*# and paste0 can be even simpler see ?paste0*

```
paste0("Visit", 1:5)
```

```
[1] "Visit1" "Visit2" "Visit3" "Visit4" "Visit5"
```

## Paste Depicting How Collapse Works

```
paste(1:5)
```

```
[1] "1" "2" "3" "4" "5"
```

```
paste(1:5, collapse = " ")
```

```
[1] "1 2 3 4 5"
```

# The `stringr` package

Like `dplyr`, the `stringr` package:

- Makes some things more intuitive
- Is different than base R
- Is used on forums for answers
- Has a standard format for most functions
  - the first argument is a string like first argument is a `data.frame` in `dplyr`



# Splitting/Find/Replace and Regular Expressions

- R can do much more than find exact matches for a whole string
- Like Perl and other languages, it can use regular expressions.
- What are regular expressions?
  - Ways to search for specific strings
  - Can be very complicated or simple
  - Highly Useful - think "Find" on steroids

## A bit on Regular Expressions

- <http://www.regular-expressions.info/reference.html>
- They can use to match a large number of strings in one statement
- `.` matches any single character
- `*` means repeat as many (even if 0) more times the last character
- `?` makes the last thing optional
- `^` matches start of vector `^a` - starts with "a"
- `$` matches end of vector `b$` - ends with "b"

# Splitting Strings

# Substringing

Very similar:

Base R

- `substr(x, start, stop)` - substrings from position start to position stop
- `strsplit(x, split)` - splits strings up - returns list!

stringr

- `str_sub(x, start, end)` - substrings from position start to position end
- `str_split(string, pattern)` - splits strings up - returns list!

## Splitting String: base R

In base R, `strsplit` splits a vector on a string into a list

```
x <- c("I really", "like writing", "R code programs")  
y <- strsplit(x, split = " ") # returns a list  
y
```

```
[[1]]  
[1] "I"      "really"
```

```
[[2]]  
[1] "like"   "writing"
```

```
[[3]]  
[1] "R"      "code"   "programs"
```

## Splitting String: **stringr**

`stringr::str_split` does the same thing:

```
library(stringr)
y2 <- str_split(x, " ") # returns a list
y2
```

```
[[1]]
[1] "I"      "really"
```

```
[[2]]
[1] "like"   "writing"
```

```
[[3]]
[1] "R"      "code"   "programs"
```

## Using a fixed expression

One example case is when you want to split on a period ".". In regular expressions . means **ANY** character, so

```
str_split("I.like.strings", ".")
```

```
[[1]]  
[1] "" "" "" "" "" "" "" "" "" "" "" "" "" "" ""
```

```
str_split("I.like.strings", fixed("."))
```

```
[[1]]  
[1] "I"      "like"   "strings"
```

## Let's extract from y

```
suppressPackageStartupMessages(library(dplyr)) # must be loaded AFTER plyr  
y[[2]]
```

```
[1] "like"    "writing"
```

```
sapply(y, dplyr::first) # on the fly
```

```
[1] "I"      "like" "R"
```

```
sapply(y, nth, 2) # on the fly
```

```
[1] "really" "writing" "code"
```

```
sapply(y, last) # on the fly
```

```
[1] "really" "writing" "programs"
```



## 'Find' functions: base R

`grep`: `grep`, `grep1`, `regexpr` and `gregexpr` search for matches to argument `pattern` within each element of a character vector: they differ in the format of and amount of detail in the results.

`grep(pattern, x, fixed=FALSE)`, where:

- `pattern` = character string containing a regular expression to be matched in the given character vector.
- `x` = a character vector where matches are sought, or an object which can be coerced by `as.character` to a character vector.
- If `fixed=TRUE`, it will do exact matching for the phrase anywhere in the vector (regular find)

## 'Find' functions: `stringr`

`str_detect`, `str_subset`, `str_replace`, and `str_replace_all` search for matches to argument `pattern` within each element of a character vector: they differ in the format of and amount of detail in the results.

- `str_detect` - returns `TRUE` if `pattern` is found
- `str_subset` - returns only the strings which `pattern` were detected
  - convenient wrapper around `x[str_detect(x, pattern)]`
- `str_extract` - returns only strings which `pattern` were detected, but ONLY the pattern
- `str_replace` - replaces `pattern` with `replacement` the first time
- `str_replace_all` - replaces `pattern` with `replacement` as many times matched

## 'Find' functions: stringr compared to base R

Base R does not use these functions. Here is a "translator" of the `stringr` function to base R functions

- `str_detect` - similar to `grep1` (return logical)
- `grep(value = FALSE)` is similar to `which(str_detect())`
- `str_subset` - similar to `grep(value = TRUE)` - return value of matched
- `str_replace` - similar to `sub` - replace one time
- `str_replace_all` - similar to `gsub` - replace many times

## Let's look at modifier for `stringr`

`?modifiers`

- `fixed` - match everything exactly
- `regexp` - default - uses **regular expressions**
- `ignore_case` is an option to not have to use `tolower`

# Important Comparisons

Base R:

- Argument order is (pattern, x)
- Uses option (`fixed = TRUE`)

`stringr`

- Argument order is (string, pattern) aka (x, pattern)
- Uses function `fixed(pattern)`

## 'Find' functions: Finding Indices

These are the indices where the pattern match occurs:

```
grep("two aliens",ufo$comments)
```

```
[1] 1730 61724
```

```
which(grepl("two aliens", ufo$comments))
```

```
[1] 1730 61724
```

```
which(str_detect(ufo$comments, "two aliens"))
```

```
[1] 1730 61724
```

## 'Find' functions: Finding Logicals

These are the indices where the pattern match occurs:

```
grep1("two aliens", ufo$comments)
```

```
[1] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[12] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[23] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[34] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[45] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[56] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[67] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[78] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[89] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[100] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[111] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[122] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[133] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[144] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[155] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[166] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[177] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[188] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
[199] FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE FALSE
```

# 'Find' functions: finding values, base R

```
grep("two aliens", ufo$comments,value=TRUE)
```

```
[1] "((HOAX??))  two aliens appeared from a bright light to peacefully investigate the surroundings in the woods"
[2] "Witnessed two aliens walking along baseball field fence."
```

```
ufo[grep("two aliens",ufo$comments),]
```

```
# A tibble: 2 x 11
  datetime          city state country    shape
  <chr>          <chr> <chr>   <chr>   <chr>
1 10/14/2006 02:00    yuma   va      us formation
2  7/1/2007 23:00 north grosvenordale ct    <NA>   unknown
# ... with 6 more variables: `duration (seconds)` <dbl>, `duration
# (hours/min)` <chr>, comments <chr>, `date posted` <chr>,
# latitude <chr>, longitude <dbl>
```



## 'Find' functions: finding values, stringr and dplyr

```
str_subset(ufo$comments, "two aliens")
```

```
[1] "((HOAX??)) two aliens appeared from a bright light to peacefully inve:  
[2] "Witnessed two aliens walking along baseball field fence."
```

```
ufo %>% filter(str_detect(comments, "two aliens"))
```

```
# A tibble: 2 x 11
```

	datetime	city	state	country	shape
	<chr>	<chr>	<chr>	<chr>	<chr>
1	10/14/2006 02:00	yuma	va	us	formation
2	7/1/2007 23:00	north grosvenordale	ct	<NA>	unknown

```
# ... with 6 more variables: `duration (seconds)` <dbl>, `duration  
# (hours/min)` <chr>, comments <chr>, `date posted` <chr>,  
# latitude <chr>, longitude <dbl>
```

## Showing difference in `str_extract`

`str_extract` extracts just the matched string

```
ss = str_extract(ufo$comments, "two aliens")  
head(ss)
```

```
[1] NA NA NA NA NA NA
```

```
ss[ !is.na(ss)]
```

```
[1] "two aliens" "two aliens"
```

# Using Regular Expressions

- Look for anycomment that starts with "aliens"

```
grep("^aliens.*", x = ufo$comments, value = TRUE)
```

```
[1] "aliens speak german???" "aliens exist"  
[3] "aliens in srilanka"
```

```
head(grep("space.?ship", x = ufo$comments, value = TRUE),3)
```

```
[1] "I saw the cylinder shaped looked like a spaceship hovering above the ea:  
[2] "description of a spaceship spotted over Birmingham Alabama in 1967."  
[3] "A space ship was descending to the ground"
```

## Using Regular Expressions: **stringr**

```
grep("^aliens.*", x = ufo$comments, value = TRUE)
```

```
[1] "aliens speak german???" "aliens exist"  
[3] "aliens in srilanka"
```

```
head(str_subset( ufo$comments, "space.?ship"),3)
```

```
[1] "I saw the cylinder shaped looked like a spaceship hovering above the ea:  
[2] "description of a spaceship spotted over Birmingham Alabama in 1967."  
[3] "A space ship was descending to the ground"
```

# Replace

Let's say we wanted to sort the data set by latitude and longitude:

```
class(ufo$latitude)
```

```
[1] "character"
```

```
sort(c("1", "2", "10")) # not sort correctly (order simply ranks the data)
```

```
[1] "1"  "10" "2"
```

```
order(c("1", "2", "10"))
```

```
[1] 1 3 2
```

# Replace

So we must change the coordinates into a numeric:

```
head(ufo$latitude, 4)
```

```
[1] "29.8830556" "29.38421"   "53.2"        "28.9783333"
```

```
head(as.numeric(ufo$latitude), 4)
```

```
Warning in head(as.numeric(ufo$latitude), 4): NAs introduced by coercion
```

```
[1] 29.88306 29.38421 53.20000 28.97833
```

## Dropping bad observations

```
dropIndex = which(is.na(as.numeric(ufo$latitude)) |  
                  is.na(as.numeric(ufo$longitude)))
```

```
Warning in which(is.na(as.numeric(ufo$latitude)) | is.na(as.numeric(ufo  
$longitude))): NAs introduced by coercion
```

```
ufo_clean = ufo[-dropIndex,]  
dim(ufo_clean)
```

```
[1] 88678    11
```

# Ordering

```
ufo2 = ufo_clean
ufo2$latitude = as.numeric(ufo2$latitude)
ufo2$longitude = as.numeric(ufo2$longitude)
ufo2 <- ufo2[order(ufo2$latitude, ufo2$longitude), ]
ufo2[1:5, c("datetime", "latitude", "longitude")]
```

```
# A tibble: 5 x 3
  datetime latitude longitude
  <chr>      <dbl>      <dbl>
1 5/15/1994 13:00 -82.86275 -135.0000
2 4/14/2002 22:22 -46.41319 168.3538
3 7/3/2002 22:35 -46.41319 168.3538
4 10/23/2008 04:45 -46.16399 169.8750
5 6/1/1998 22:00 -45.74676 170.5691
```



## Replacing and subbing: **stringr**

We can do the same thing (with 2 piping operations!) in dplyr

```
ufo_dplyr = ufo_clean
ufo_dplyr = ufo_dplyr %>% mutate(
  latitude = latitude %>% as.numeric,
  longitude = longitude %>% as.numeric) %>%
  arrange(latitude, longitude)
ufo_dplyr[1:5, c("datetime", "latitude", "longitude")]
```

```
# A tibble: 5 x 3
  datetime latitude longitude
  <chr>      <dbl>      <dbl>
1 5/15/1994 13:00 -82.86275 -135.0000
2 4/14/2002 22:22 -46.41319 168.3538
3 7/3/2002 22:35 -46.41319 168.3538
4 10/23/2008 04:45 -46.16399 169.8750
5 6/1/1998 22:00 -45.74676 170.5691
```